

A Unifying Class of Algorithms for Semi-Streaming Bipartite Maximum Matching

Kheeran Naidu

University of Bristol

kn16063@bristol.ac.uk

Joint work with Dr Christian Konrad

Overview

1 Background

2 Our Work

- Algorithmic
- Impossibility

3 Discussion

Matchings


Definition

A **(bipartite) matching** is a subset of edges of a graph where every vertex has degree at most 1.

Matchings

Definition


A **(bipartite) matching** is a subset of edges of a graph where every vertex has degree at most 1.



Matchings

Definition

A **(bipartite) matching** is a subset of edges of a graph where every vertex has degree at most 1.




matching


Matchings

Definition

A **(bipartite) matching** is a subset of edges of a graph where every vertex has degree at most 1.



matching




maximal matching


Matchings

Definition


A **(bipartite) matching** is a subset of edges of a graph where every vertex has degree at most 1.



matching



maximal matching




maximum matching


Matchings

Definition


A **(bipartite) matching** is a subset of edges of a graph where every vertex has degree at most 1.



matching




maximal matching



optimal matching

Esfandiari et al. [ICDMW16]

A **semi-incomplete matching** is an extended matching in bipartite graphs where the vertices of one bipartition allows for degree $d > 1$.




semi-incomplete matching


Matchings

Definition


A **(bipartite) matching** is a subset of edges of a graph where every vertex has degree at most 1.



matching



maximal matching



optimal matching


Matchings

Definition

A **(bipartite) matching** is a subset of edges of a graph where every vertex has degree at most 1.

Finding a maximum (optimal) matching:

- exact algorithms exists, i.e. Hopcroft-Karp [SWAT71];
- require **random access** to the graph's edges (infeasible requirement for massive graphs).




Semi-Streaming Model

Feigenbaum et al. [ICALP04]

A graph with n vertices is presented to an algorithm as a stream of edges where the storage space of the algorithm is bounded by $O(n \text{ polylog } n)$.

- Only allows **sequential access** to the graph.
- Algorithms with space $O(n \text{ polylog } n) = O(n (\log n)^{O(1)})$.
- Ideally with few passes of the stream.



Two-Pass Bipartite Maximum (Optimal) Matching

Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).

Two-Pass Bipartite Maximum (Optimal) Matching

Class of algorithms:

- ① finds a maximal matching M (first-pass);
- ② increases the size of the matching M (second-pass).

Note: Only strategy proven to work.

Two-Pass Bipartite Maximum (Optimal) Matching

Algorithmic:

Impossibility:

Two-Pass Bipartite Maximum (Optimal) Matching

Algorithmic:

- unify the two dominant techniques; vertex subsampling and finding a semi-incomplete matching;

Impossibility:

Two-Pass Bipartite Maximum (Optimal) Matching

Algorithmic:

- unify the two dominant techniques; vertex subsampling and finding a semi-incomplete matching;
- present a wider set of algorithms which contains the most recent state-of-the-art results;

Impossibility:

Two-Pass Bipartite Maximum (Optimal) Matching

Algorithmic:

- unify the two dominant techniques; vertex subsampling and finding a semi-incomplete matching;
- present a wider set of algorithms which contains the most recent state-of-the-art results;
- find a novel meta algorithm that exactly achieves the current state-of-the-art $2 - \sqrt{2} \approx \frac{1}{2} + 0.085$ -approximation.

Impossibility:

Two-Pass Bipartite Maximum (Optimal) Matching

Algorithmic:

- unify the two dominant techniques; vertex subsampling and finding a semi-incomplete matching;
- present a wider set of algorithms which contains the most recent state-of-the-art results;
- find a novel meta algorithm that exactly achieves the current state-of-the-art $2 - \sqrt{2} \approx \frac{1}{2} + 0.085$ -approximation.

Impossibility:

- In this class of algorithms, (*the first pass finds only a maximal matching*), no better than a $\frac{2}{3} \approx \frac{1}{2} + 0.167$ -approximation is possible in the semi-streaming model.

	Algorithmic	Impossibility
one-pass		$\frac{1}{2} + 0.167$ [SODA12]
	$\frac{1}{2}$ [ICALP04]	$\frac{1}{2} + 0.132$ [SODA13]
		$\frac{1}{2} + 0.091$ [SODA21]
two-pass	$\frac{1}{2} + 0.019$ [APPROX12]	
	$\frac{1}{2} + 0.083$ [ICDMW16]	
	$\frac{1}{2} + 0.063$ [APPROX17]	$\frac{1}{2} + 0.167^1$
	$\frac{1}{2} + 0.085$ [MFCS18]	
	$\frac{1}{2} + 0.085$	

¹where the first pass finds a maximal matching, i.e., at least a $\frac{1}{2}$ -approximation.

Overview

1 Background

2 Our Work

- Algorithmic
- Impossibility

3 Discussion

Overview

1 Background

2 Our Work

- Algorithmic
- Impossibility

3 Discussion

Algorithmic Approach

Class of algorithms:

- ① finds a maximal matching M (first-pass);
- ② increases the size of the matching M (second-pass).

Algorithmic Approach

Class of algorithms:

- ① finds a maximal matching M (first-pass);
- ② increases the size of the matching M (second-pass).

Definition

An **augmenting path** begins and ends with an edge not in the matching, and alternates along the path between edges in and out of the matching.


Algorithmic Approach

Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).

Definition

An **augmenting path** begins and ends with an edge not in the matching, and alternates along the path between edges in and out of the matching.




Algorithmic Approach

Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).

Definition

An **augmenting path** begins and ends with an edge not in the matching, and alternates along the path between edges in and out of the matching.




Algorithmic Approach

Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).

Definition

An **augmenting path** begins and ends with an edge not in the matching, and alternates along the path between edges in and out of the matching.




Algorithmic Approach

Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).

Definition

An **augmenting path** begins and ends with an edge not in the matching, and alternates along the path between edges in and out of the matching.




Algorithmic Approach

Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).

Definition

An **augmenting path** begins and ends with an edge not in the matching, and alternates along the path between edges in and out of the matching.




Algorithmic Approach

Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).

Definition

An **augmenting path** begins and ends with an edge not in the matching, and alternates along the path between edges in and out of the matching.




Algorithmic Approach

Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).

Definition


An **augmenting path** begins and ends with an edge not in the matching, and alternates along the path between edges in and out of the matching.



Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

The goal of the algorithm is to find the maximum (optimal) matching M^* .




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$




Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.

Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$



Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance

(worst)

Let π_G

edges


order.

Using the two dominant techniques:

- 1 subsample the inner vertices [APPROX12] [MFCS18]
- 2 run GREEDY_d [ICDMW16] [APPROX17]

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$



M

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance

(worst)

Let π

edges


order.

First

• Λ

Using the two dominant techniques:

- 1 subsample the inner vertices [APPROX12] [MFCS18]
- 2 run GREEDY_d [ICDMW16] [APPROX17]



semi-incomplete matching

IVI

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



$$M'$$
$$d = 3, p = 0.67$$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



$$M'$$
$$d = 3, p = 0.67$$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



$$M'$$
$$d = 3, p = 0.67$$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



M'

$d = 3, p = 0.67$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



$$d = 3, p = 0.67$$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



M'

$d = 3, p = 0.67$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



$$M'$$
$$d = 3, p = 0.67$$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



$$d = 3, p = 0.67$$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



M'

$d = 3, p = 0.67$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:

- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



$$d = 3, p = 0.67$$

$$\mathbb{E}_{M'}[|M'_L|] = \frac{dp}{d+p} \cdot |M_L^*|$$

Class of Algorithms

Let $G = (A, B, E)$ be a hard-instance (worst-case) graph.


Let $\pi_G = e_1, e_2, \dots$ be a stream of its edges in adversarial (worst-case) order.

First pass:

- $M \leftarrow \text{GREEDY}(\pi_G)$

Second pass:


- subsample $M' \subseteq M$ with prob. p
- $S_L \leftarrow E \cap A_{out} \times B(M')$
- $S_R \leftarrow E \cap A(M') \times B_{out}$
- $M'_L \leftarrow \text{GREEDY}_d(S_L \cap \pi_G)$
- $M'_R \leftarrow \text{GREEDY}_d(S_R \cap \pi_G)$



$$d = 3, p = 0.67$$

$$\mathbb{E}_{M'}[|M'_L|] \geq \frac{dp}{d+p} \cdot |M_L^*|$$


Main Theorem (Proof Outline)



Main Theorem (Proof Outline)

Setup:


- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;



Main Theorem (Proof Outline)

Setup:


- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;



Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?




Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?



$$\mathbb{E}[|M_{B'}^*|] \leq (1 + \frac{p}{d})\mathbb{E}[|M|]$$


Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?



$$p \cdot |M^*| \leq (1 + \frac{p}{d}) \mathbb{E}[|M|]$$


Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.

Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?



$$\mathbb{E}[|M|] \geq \frac{dp}{d+p} \cdot |M^*|$$


Main Theorem (Proof Outline)

Setup:

- any bipartite graph $G = (A, B, E)$ with a maximum (optimal) matching M^* ;
- subsample $B' \subseteq B$ with prob. p ;
- $M \leftarrow \text{GREEDY}_d$.


Proof:

- $M_{B'}^* \leftarrow \{ab \in M^* : b \in B'\}$;
- For every edge that is added by GREEDY_d , how many edges in $M_{B'}^*$ are blocked?
- Formalised using Wald's Equation.



$$\mathbb{E}[|M|] \geq \frac{dp}{d+p} \cdot |M^*|$$


Optimal Algorithms I



$$d = 3, p = 0.67$$

$$\mathbb{E}_{M'}[|M'_L|] = \frac{dp}{d+p} \cdot |M_L^*|$$

Optimal Algorithms I




$$d = 3, p = 0.67$$

$$\mathbb{E}_{M'}[|M'_L|] = \frac{dp}{d+p} \cdot |M_L^*|$$

Optimal Algorithms I

Analysis:


We find at least $(\frac{1}{2} + \frac{p}{d+p} - \frac{p}{2d}) \cdot |M^*|$ edges in our final matching.




$$d = 3, p = 0.67$$

$$\mathbb{E}_{M'}[|M'_L|] = \frac{dp}{d+p} \cdot |M^*|$$


Optimal Algorithms II




Optimal Algorithms II




Optimal Algorithms II



Optimal Algorithms II



Optimal Algorithms II



Overview


1 Background

2 Our Work

- Algorithmic
- Impossibility


3 Discussion

Two-Party Communication Setup



Proving Impossibility Results

Two-Party Communication Setup




Goal:

- to bound the size of the message from Alice to Bob needed to output a large matching.

One-Pass Impossibility Proof [SODA12]

Outline:



- 1 A family of graphs constructed from very dense Rusza-Szemerédi (RS) graphs.



One-Pass Impossibility Proof [SODA12]

Outline:



- 1 A family of graphs constructed from very dense Rusza-Szemerédi (RS) graphs.



One-Pass Impossibility Proof [SODA12]

Outline:



- 1 A family of graphs constructed from very dense Rusza-Szemerédi (RS) graphs.



One-Pass Impossibility Proof [SODA12]

Outline:



- 1 A family of graphs constructed from very dense Rusza-Szemerédi (RS) graphs.



One-Pass Impossibility Proof [SODA12]

Outline:



- 1 A family of graphs constructed from very dense Rusza-Szemerédi (RS) graphs.
- 2 Prove that a $(\frac{2}{3} + \epsilon)$ -approx requires $N^{1+\Omega(\frac{1}{\log \log N})} \supset O(N \text{ polylog } N)$ space.



Two-Pass Impossibility Proof



Class of algorithms:

- 1 finds a maximal matching M (first-pass);
- 2 increases the size of the matching M (second-pass).



Two-Pass Impossibility Proof

Goal: Give both Alice and Bob knowledge of a maximal matching without affecting the difficulty of the problem.





Two-Pass Impossibility Proof

Goal: Give both Alice and Bob knowledge of a maximal matching without affecting the difficulty of the problem.

Outline:

- We extend [SODA12]'s construction a family of RS graphs.





Two-Pass Impossibility Proof

Goal: Give both Alice and Bob knowledge of a maximal matching without affecting the difficulty of the problem.

Outline:

- We extend [SODA12]'s construction a family of RS graphs.
- Do dense RS graphs contain perfect matchings?





Two-Pass Impossibility Proof

Goal: Give both Alice and Bob knowledge of a maximal matching without affecting the difficulty of the problem.

Outline:

- We extend [SODA12]'s construction a family of RS graphs.
- Do dense RS graphs contain perfect matchings?
- Find a near-perfect matching of size $N - \epsilon'N$





Two-Pass Impossibility Proof

Goal: Give both Alice and Bob knowledge of a maximal matching without affecting the difficulty of the problem.

Outline:

- We extend [SODA12]'s construction a family of RS graphs.
- Do dense RS graphs contain perfect matchings?
- Find a near-perfect matching of size $N - \epsilon' N$





Two-Pass Impossibility Proof

Goal: Give both Alice and Bob knowledge of a maximal matching without affecting the difficulty of the problem.

Outline:

- We extend [SODA12]'s construction a family of RS graphs.
- Do dense RS graphs contain perfect matchings?
- Find a near-perfect matching of size $N - \epsilon' N$





Two-Pass Impossibility Proof

Goal: Give both Alice and Bob knowledge of a maximal matching without affecting the difficulty of the problem.

Outline:

- We extend [SODA12]'s construction a family of RS graphs.
- Do dense RS graphs contain perfect matchings?
- Find a near-perfect matching of size $N - \epsilon' N$
- $(\frac{2}{3} + \epsilon)$ -approx requires space $N^{1+\Omega(\frac{1}{\log \log N})} \supset O(N \text{ polylog } N)$.



Overview

1 Background

2 Our Work


- Algorithmic
- Impossibility

3 Discussion

Conclusion


Conclusion

- Our set algorithms unifies the dominant techniques used to tackle the two-pass bipartite maximum matching problem, achieving the current state-of-the-art.




Conclusion

- Our set algorithms unifies the dominant techniques used to tackle the two-pass bipartite maximum matching problem, achieving the current state-of-the-art.
- For appropriate settings of d and p , we can find Konrad's [MFCS18] and Esfandiari et al.'s [ICDMW16] algorithms.




Conclusion

- Our set algorithms unifies the dominant techniques used to tackle the two-pass bipartite maximum matching problem, achieving the current state-of-the-art.
- For appropriate settings of d and p , we can find Konrad's [MFCS18] and Esfandiari et al.'s [ICDMW16] algorithms.




Conclusion

- Our set algorithms unifies the dominant techniques used to tackle the two-pass bipartite maximum matching problem, achieving the current state-of-the-art.
- For appropriate settings of d and p , we can find Konrad's [MFCS18] and Esfandiari et al.'s [ICDMW16] algorithms.




Conclusion

- Our set algorithms unifies the dominant techniques used to tackle the two-pass bipartite maximum matching problem, achieving the current state-of-the-art.
- For appropriate settings of d and p , we can find Konrad's [MFCS18] and Esfandiar et al.'s [ICDMW16] algorithms.
- Our hard-instance graph proves that our analysis is tight.



Conclusion

- Our set algorithms unifies the dominant techniques used to tackle the two-pass bipartite maximum matching problem, achieving the current state-of-the-art.
- For appropriate settings of d and p , we can find Konrad's [MFCS18] and Esfandiar et al.'s [ICDMW16] algorithms.
- Our hard-instance graph proves that our analysis is tight.
- We reduced the gap of possibility with this class of algorithms to $[0.585, 0.667]$.



Open Questions

- Can we extend other one-pass impossibility results to improve the two-pass bound? I.e. Kapralov's [SODA21].
- Is there a way to do better by finding more than just a maximal matching in the first-pass?
- Can we beat a $\frac{1}{2}$ -approximation in just one-pass?

Thank You